

Networks and Cloud Infrastructures

Università degli Studi di Napoli Federico II

Dipartimento di Ingegneria Elettrica e delle Tecnologie dell'Informazione DIETI

Laurea Magistrale in Ingegneria Informatica

Prof. Giorgio Ventre

The Mininet network emulator



- ▶ Mininet is a lightweight virtualization/container based emulator that allows to reproduce virtual networks with arbitrary topologies
 - ▶ modest hardware requirements, fast startup, hundreds of nodes
 - ▶ command line tool, CLI, simple Python API
- ▶ Mininet enables SDN development on any laptop or PC, and SDN designs can move seamlessly between Mininet (allowing inexpensive and streamlined development), and the real hardware running at line rate in live deployments
- ▶ Python scripts can be used to orchestrate an experiment
 - ▶ Network topology definition
 - ▶ Events to be triggered in network nodes (e.g. execution of a program)
- ▶ Mininet VM installation: the easiest way of installing Mininet
 1. Download and install on your PC one of the available hypervisors (e.g., VirtualBox, or KVM)
 2. Download the Mininet pre-installed VM image
 3. Import VM image into selected hypervisor

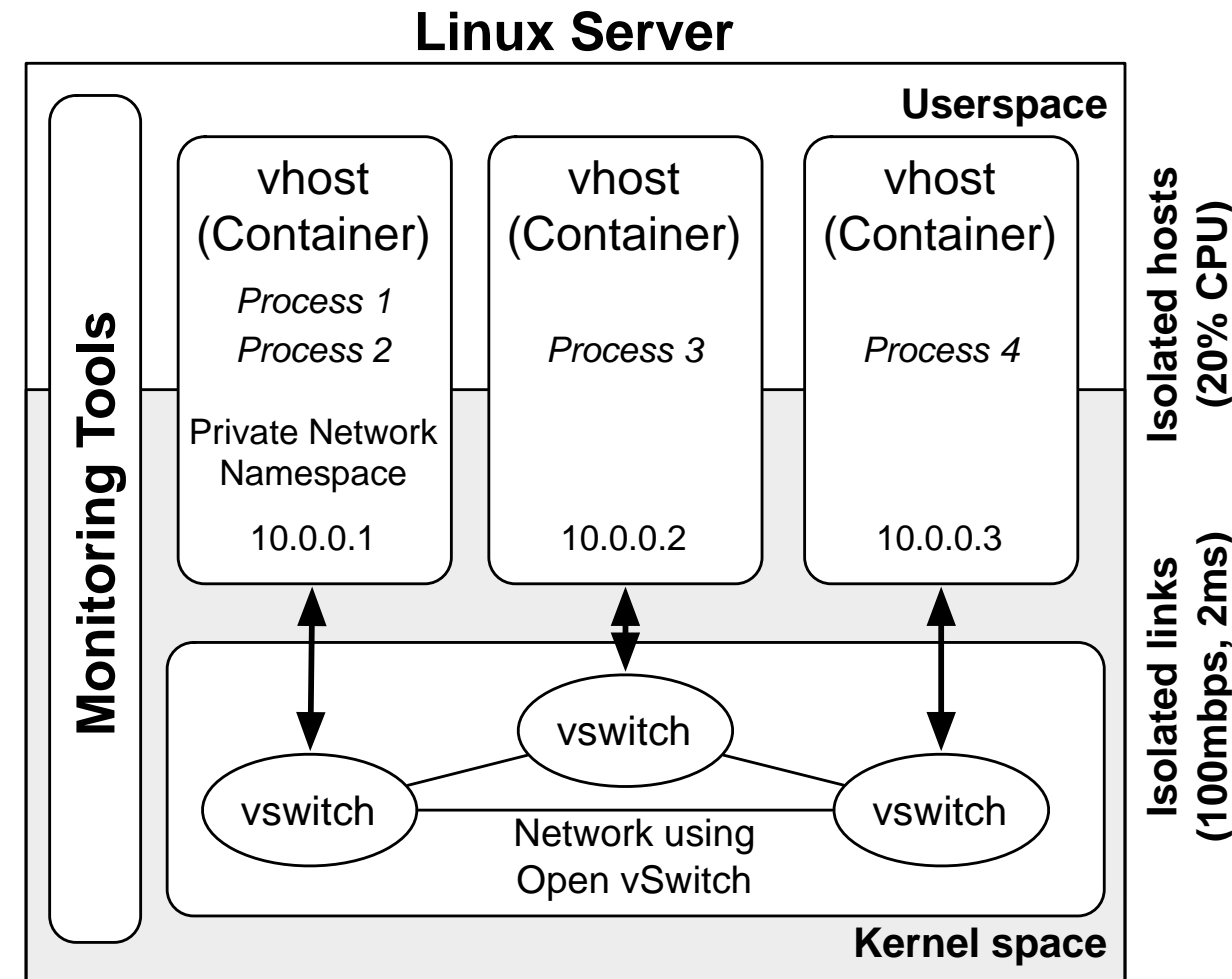


- ▶ <http://mininet.org/walkthrough/>
- ▶ <https://github.com/mininet/mininet/wiki/Introduction-to-Mininet>
- ▶ <https://github.com/mininet/mininet/wiki/Videos>
- ▶ <https://github.com/mininet/mininet/wiki/Documentation>
- ▶ Floodlight Controller
<https://floodlight.atlassian.net/wiki/display/floodlightcontroller/Getting+Started>
- ▶ RYU Controller Tutorial
<http://sdnhub.org/tutorials/ryu/>

Mininet: emulated hosts and switches



- ▶ Mininet combines lightweight virtualization (*containers*) with software switches to emulate networks in a Linux-based system
- ▶ A Mininet network consists of:
 - ▶ **ISOLATED HOSTS:** a group of user-level processes moved into a network namespace that provides exclusive ownership of interfaces, ports and routing tables
 - ▶ **EMULATED LINKS:** Linux Traffic Control (tc) enforces the data rate of each link to shape traffic to a configured rate - each emulated host has its own virtual Ethernet interface(s)
 - ▶ **EMULATED SWITCHES:** the default Linux Bridge or the Open vSwitch running in kernel mode are used to switch packets across interfaces
- ▶ Emulated hosts share the same Linux kernel and the file system of the host in which they run



Prerequisites



- ▶ Install Virtualbox, VMWare or other, and create an Ubuntu 22.04 VM
- ▶ Install the following packages on the VM using the provided command line:

```
$ sudo apt install git
$ sudo apt install python3-pip
$ sudo pip3 install pandas
$ pip install ryu
$ sudo apt install d-itg
$ sudo apt install nload
```

- ▶ Clone Mininet Repository and Install:

```
$ git clone https://github.com/mininet/mininet
$ cd mininet
$ git tag
$ git checkout -b mininet-2.3.0 2.3.0
$ cd ..
$ sudo PYTHON=python3 mininet/util/install.sh -nv
$ sudo mn --switch ovsbr --test pingall # Test Mininet installation
```

Mininet: network topologies



- ▶ Start mininet with a minimal topology (1 switch and 2 connected hosts)

```
$ sudo mn
```

is equivalent to:

```
$ sudo mn --topo minimal
```

- ▶ Start mininet with 1 switch and n connected hosts

```
$ sudo mn --topo single,n
```

- ▶ Start mininet with a linear topology (n switches in a row and 1 host connected to each switch)

```
$ sudo mn --topo linear,n
```

- ▶ Start mininet with a tree topology with depth n and fanout m

```
$ sudo mn --topo tree,depth=n,fanout=m
```

- ▶ Start mininet with a custom topology *mytopo* defined in a Python script (mytopo.py)

```
$ sudo mn --custom mytopo.py --topo mytopo
```

- ▶ Stop mininet processes

```
$ sudo mn -c
```

```
class MyTopo( Topo ):
    def build( self, ...):
def myTest( net ):
    ...
topos = { 'mytopo': MyTopo }
tests = { 'mytest': myTest }
```

Mininet: controller option



- ▶ Start a minimal topology with the default internal controller

```
$ sudo mn
```

- ▶ Start a minimal topology without a controller

```
$ sudo mn --controller none
```

- ▶ Start a minimal topology using the reference OpenFlow controller

```
$ sudo mn --controller ref
```

- ▶ Start a minimal topology using an external controller (e.g. Ryu, Floodlight, etc.)

```
$ sudo mn --controller remote,ip=[IP_ADDDR],port=[listening_port]
```

- ▶ Start a minimal topology using an external controller on 127.0.0.1:6653

```
$ sudo mn --controller remote
```

- ▶ Start mininet by assigning MAC addresses sequentially to hosts

```
$ sudo mn --mac
```

- ▶ E.g. host h1 gets MAC 00:00:00:00:00:01, etc.

- ▶ Start mininet and show an xterm for every host and switch

```
$ sudo mn -x
```

- ▶ Start mininet and run test function iperf when the whole node is up

```
$ sudo mn --test iperf
```

or

```
# iperf -s
```

```
# iperf -c ip_address
```

- ▶ Start a minimal topology and exits: computes the time to bring the network up

```
$ sudo mn --test none
```

- ▶ Use Open vSwitch for network nodes

```
$ sudo mn --switch ovsk
```

- ▶ Force Open vSwitch to use OpenFlow protocol version OpenFlow1.3

```
$ sudo mn --switch ovsk,protocols=OpenFlow13
```


Mininet: other options



- ▶ Start mininet by assigning specific parameters to all links

```
$ sudo mn --link tc,bw=[bandwidth],delay=[delay_in_millisecond]
```

- ▶ Assigns a given bandwidth and delay to links

- ▶ Example (10.00Mbit 10ms delay):

```
$ sudo mn --link tc,bw=10,delay=10ms
```

- ▶ Open a terminal on both hosts and run iperf

```
mininet> xterm h1
```

```
mininet> xterm h2
```



- ▶ Display nodes

```
mininet> nodes
```

- ▶ Display links

```
mininet> net
```

- ▶ Dump information about all nodes

```
mininet> dump
```

- ▶ Execute a method through invoking mininet API

```
mininet> py [mininet_name_space].[method]
```



- ▶ Check the IP address of a certain node

```
mininet> h1 ifconfig -a
```

- ▶ Print the process list from a host process

```
mininet> h1 ps -a
```

- ▶ Verify the connectivity by pinging from host h1 to host h2

```
mininet> h1 ping -c 1 h2
```

- ▶ Verify the connectivity between all pairs of hosts

```
mininet> pingall
```

- ▶ Measure end-to-end bandwidth between two hosts with iperf

- ▶ Server endpoint

```
mininet> iperf -s -u -p [port_num] &
```

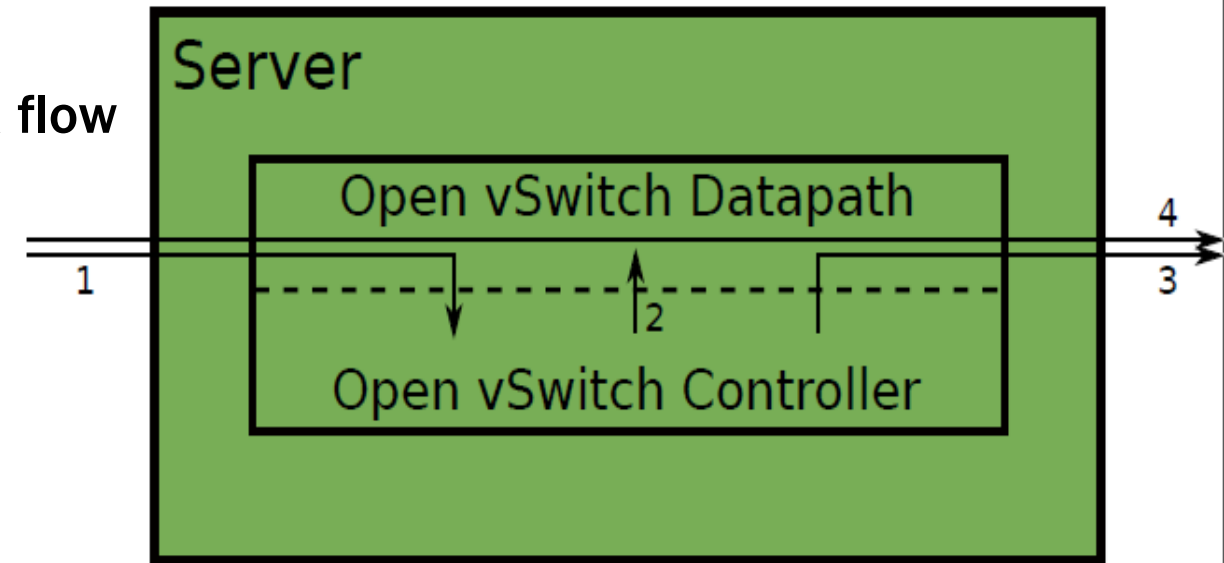
- ▶ Client endpoint

```
mininet> iperf -c [IP] -u -t [duration] -b [bandwidth] -p [port_num]
```

Open vSwitch



- ▶ Mininet is often used to instantiate networks of Open vSwitch switches
- ▶ Open vSwitch is a production quality, multilayer virtual switch licensed under the open source Apache 2.0 license
- ▶ Open vSwitch design choices:
 - ▶ Flexible Controller computation in User space
 - ▶ Fast Datapath packet handling in Kernel space
- ▶ The 1st packet of a flow is sent to the controller
- ▶ The controller programs the datapath's actions for a flow
 - ▶ Usually one, but may be a list
- ▶ Actions include:
 - ▶ Forward to a port or ports
 - ▶ Mirror
 - ▶ Encapsulate and forward to controller
 - ▶ Drop
- ▶ And it returns the packet to the datapath
 - ▶ Subsequent packets are handled directly by the datapath
- ▶ An Open vSwitch can be remotely controlled by an Openflow SDN controller



Open vSwitch commands: ovs-vsctl



Ovs:-vsctl:utility for querying and configuring ovs-vswitchd

```
$ sudo ovs-vsctl --help #
```

```
usage: ovs-vsctl [OPTIONS] COMMAND [ARG...]
```

Open vSwitch commands:

init	initialize database, if not yet initialized
show	print overview of database contents
emer-reset	reset configuration to clean state

Bridge commands:

add-br BRIDGE	create a new bridge named BRIDGE
add-br BRIDGE PARENT VLAN	create new fake BRIDGE in PARENT on VLAN
del-br BRIDGE	delete BRIDGE and all of its ports
list-br	print the names of all the bridges

...

Port commands (a bond is considered to be a single port):

list-ports BRIDGE	print the names of all the ports on BRIDGE
add-port BRIDGE PORT	add network device PORT to BRIDGE

...

del-port [BRIDGE] PORT	delete PORT (which may be bonded) from BRIDGE
------------------------	---

Open vSwitch commands: ovs-vsctl (cont.)



...

Interface commands (a bond consists of multiple interfaces):

<code>list-ifaces BRIDGE</code>	print the names of all interfaces on BRIDGE
<code>iface-to-br IFACE</code>	print name of bridge that contains IFACE

Controller commands:

<code>get-controller BRIDGE</code>	print the controllers for BRIDGE
<code>del-controller BRIDGE</code>	delete the controllers for BRIDGE
<code>[--inactivity-probe=MSECS]</code>	
<code>set-controller BRIDGE TARGET...</code>	set the controllers for BRIDGE
<code>get-fail-mode BRIDGE</code>	print the fail-mode for BRIDGE
<code>del-fail-mode BRIDGE</code>	delete the fail-mode for BRIDGE
<code>set-fail-mode BRIDGE MODE</code>	set the fail-mode for BRIDGE to MODE

Manager commands:

<code>get-manager</code>	print the managers
<code>del-manager</code>	delete the managers

...

Open vSwitch commands: ovs-ofctl



Ovs:-ofctl: command line tool for monitoring and administering OpenFlow switches

```
$ sudo ovs-ofctl --help
```

```
ovs-ofctl: OpenFlow switch management utility
```

```
usage: ovs-ofctl [OPTIONS] COMMAND [ARG...]
```

For OpenFlow switches:

```
show SWITCH
```

show OpenFlow information

```
dump-desc SWITCH
```

print switch description

```
dump-tables SWITCH
```

print table stats

```
...
```

```
dump-ports SWITCH [PORT]
```

print port statistics

```
dump-ports-desc SWITCH
```

print port descriptions

```
dump-flows SWITCH
```

print all flow entries

```
dump-flows SWITCH FLOW
```

print matching FLOWS

```
dump-aggregate SWITCH
```

print aggregate flow statistics

```
dump-aggregate SWITCH FLOW
```

print aggregate stats for FLOWS

```
add-flow SWITCH FLOW
```

add flow described by FLOW

```
add-flows SWITCH FILE
```

add flows from FILE

```
mod-flows SWITCH FLOW
```

modify actions of matching FLOWS

```
del-flows SWITCH [FLOW]
```

delete matching FLOWS

Experiment #1: internal controller



- ▶ Start Mininet and create a simple topology with 1 switch and 2 hosts

```
$ sudo mn --topo single,2 --mac --switch ovsk
```

- ▶ By default, Mininet creates an internal controller that implements a simple learning switch functionality

- ▶ Hosts are named h1 and h2

- ▶ Open xterms for hosts h1 and h2 from Mininet prompt

```
mininet> xterm h1 h2
```

- ▶ Open wireshark from xterm on h1

```
h1# sudo wireshark
```

- ▶ Let h1 ping h2 from Mininet prompt

```
mininet> h1 ping h2
```

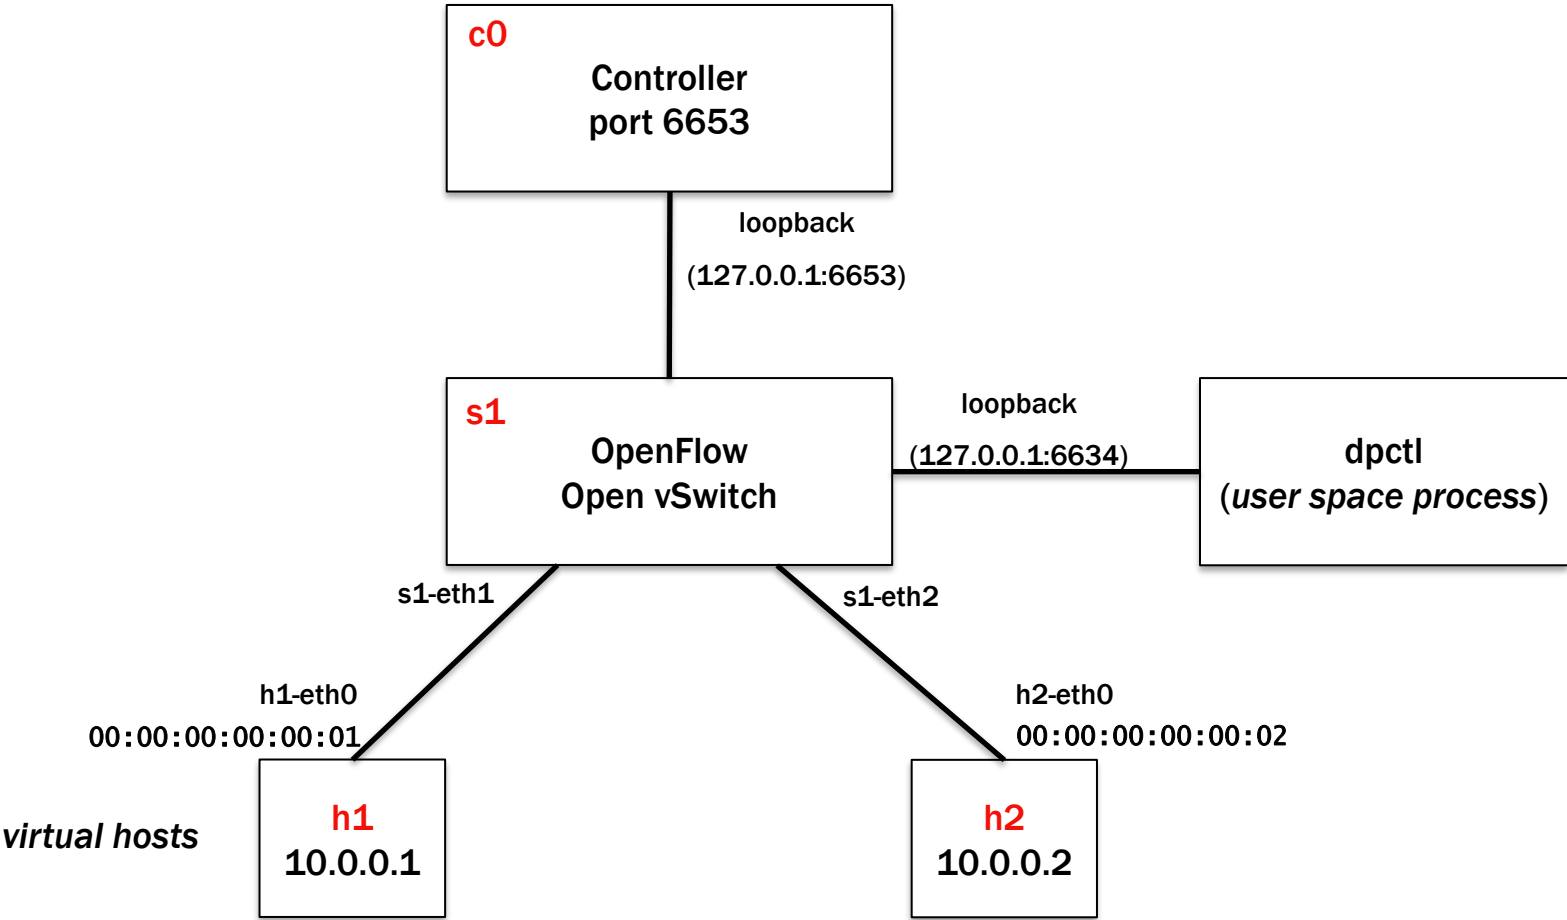
- ▶ Exec a simple web server (listening on port 80) from xterm on h2

```
h2# python3 -m http.server 80
```

- ▶ Let h1 ping h2 from Mininet prompt

```
mininet> h1 wget h2
```


Experiment #1: components



Open vSwitch CLI commands (1)



▶ **sudo ovs-vsctl show**

Lists all instances of Open vSwitch (e.g. s1, s2, ...)

```
8786cd90-73f6-43c9-bafe-72d80af2a23a
```

```
Bridge s1
```

```
Controller "ptcp:6654"
```

```
fail_mode: secure
```

```
Port s1
```

```
Interface s1
```

```
type: internal
```

```
Port s1-eth1
```

```
Interface s1-eth1
```

```
Port s1-eth2
```

```
Interface s1-eth2
```

```
ovs_version: "2.13.0"
```

▶ **sudo ovs-ofctl show s1**

Lists all ports of an Open vSwitch switch

```
OFPT_FEATURES_REPLY (xid=0x2): dpid:0000000000000001
```

```
n_tables:254, n_buffers:0
```

```
capabilities: FLOW_STATS TABLE_STATS PORT_STATS QUEUE_STATS ARP_MATCH_IP
```

```
actions: output enqueue set_vlan_vid set_vlan_pcp strip_vlan mod_dl_src mod_dl_dst  
mod_nw_src mod_nw_dst mod_nw_tos mod_tp_src mod_tp_dst
```

```
1(s1-eth1): addr:2a:de:31:4d:51:b9
```

```
config: 0
```

```
state: 0
```

```
current: 10GB-FD COPPER
```

```
speed: 10000 Mbps now, 0 Mbps max
```

```
2(s1-eth2): addr:7e:76:4f:50:13:c6
```

```
config: 0
```

```
state: 0
```

```
current: 10GB-FD COPPER
```

```
speed: 10000 Mbps now, 0 Mbps max
```

```
LOCAL(s1): addr:42:c4:0e:3a:9e:45
```

```
config: PORT_DOWN
```

```
state: LINK_DOWN
```

```
speed: 0 Mbps now, 0 Mbps max
```

```
OFPT_GET_CONFIG_REPLY (xid=0x4): frags=normal miss_send_len=0
```

Open vSwitch CLI commands (2)



▶ `sudo ovs-ofctl dump-flows s1 [-O OpenFlow13]`

Lists all entries in a switch Flow Tables; by default talks OpenFlow 1.0

▶ `sudo ovs-ofctl add-flow s1 in_port=1,actions=output:2 [-O OpenFlow13]`

Adds a flow entry into switch s1

Experiment #2: no controller



```
$ sudo mn --topo single,2 --mac --switch ovsk --controller none
```

```
$ sudo ovs-ofctl show s1 -O OpenFlow13
```

```
$ sudo ovs-ofctl dump-flows s1 -O OpenFlow13
```

```
mininet> h1 ping h2
```

All ports of switch shown, but no flows installed.
Ping fails because ARP cannot go through

```
$ sudo ovs-ofctl add-flow s1 -O OpenFlow13  
in_port=1,actions=output:2
```

```
$ sudo ovs-ofctl add-flow s1 -O OpenFlow13  
in_port=2,actions=output:1
```

```
mininet> h1 ping h2
```

```
mininet> h1 ping h2
```

Ping works now!

OpenFlow rules set by an external controller



- ▶ For the topology created with

```
$ sudo mn --topo single,2 --mac --switch ovsk --controller remote,port=6653
```

- ▶ if the OpenFlow controller is running (e.g, Floodlight with the Forwarding module enabled)

```
$ sudo ovs-ofctl dump-flows s1 -O OpenFlow13
```

- ▶ produces a similar output (actual flow rules depend on the controller logic):

```
cookie=0x2000000000000000, duration=1.343s, table=0, n_packets=0, n_bytes=0, idle_timeout=5, idle_age=1, priority=1,arp,in_port=1,dl_src=00:00:00:00:00:01,dl_dst=00:00:00:00:00:02 actions=output:2
```

```
cookie=0x2000000000000000, duration=6.361s, table=0, n_packets=1, n_bytes=42, idle_timeout=5, idle_age=1, priority=1,arp,in_port=2,dl_src=00:00:00:00:00:02,dl_dst=00:00:00:00:00:01 actions=output:1
```

```
cookie=0x2000000000000000, duration=6.356s, table=0, n_packets=6, n_bytes=588, idle_timeout=5, idle_age=0, priority=1,ip,in_port=1,dl_src=00:00:00:00:00:01,dl_dst=00:00:00:00:00:02,nw_src=10.0.0.1,nw_dst=10.0.0.2 actions=output:2
```

```
cookie=0x2000000000000000, duration=6.354s, table=0, n_packets=6, n_bytes=588, idle_timeout=5, idle_age=0,priority=1,ip,in_port=2,dl_src=00:00:00:00:00:02,dl_dst=00:00:00:00:00:01,nw_src=10.0.0.2, nw_dst=10.0.0.1 actions=output:1
```

Running Mininet scenarios described in Python scripts



- ▶ Mininet installation makes all Mininet classes available to the Python interpreter so that they can be instantiated from a regular Python script
- ▶ Instead of running the `mn` command and issuing commands from the CLI, a Mininet scenario may be executed by passing a Python script to the regular Python interpreter
- ▶ The script needs super-user rights, hence it must be run with `sudo`

```
$ sudo python test1.py
```

- ▶ The Python script needs to import all the relevant classes, create a custom topology (as a class derived from the `mininet.topo.Topo` base class), instantiate a Mininet network object and perform all the required actions
- ▶ Before terminating the script, some cleanup methods need to be invoked
- ▶ Before presenting some sample scripts, a first look at the hierarchy of classes created by Mininet is necessary
- ▶ The Python script may build the emulated scenarios by working at different semantic levels

- ▶ `mininet.node.Node`
 - ▶ A virtual network node, which is simply in a network namespace
 - ▶ From the generic Node class three classes are derived: Host, Switch and Controller
- ▶ `mininet.link.Link`
 - ▶ A basic link, which is represented as a pair of nodes

```
h1 = Host( 'h1' )
h2 = Host( 'h2' )
s1 = OVSSwitch( 's1', inNamespace=False )
c0 = Controller( 'c0', inNamespace=False )
Link( h1, s1 )
Link( h2, s1 )
h1.setIP( '10.0.0.1/8' )
h2.setIP( '10.0.0.2/8' )
```

Class	Method	Description
Node	MAC/setMAC	Return/Assign MAC address of a node or specific interface
	IP/setIP	Return/Assign IP address of a node or specific interface
	cmd	Send a command, wait for output, and return it
	terminate	Send kill signal to Node and clean up after it
Link	Link	Create a link to another node, make two new interfaces

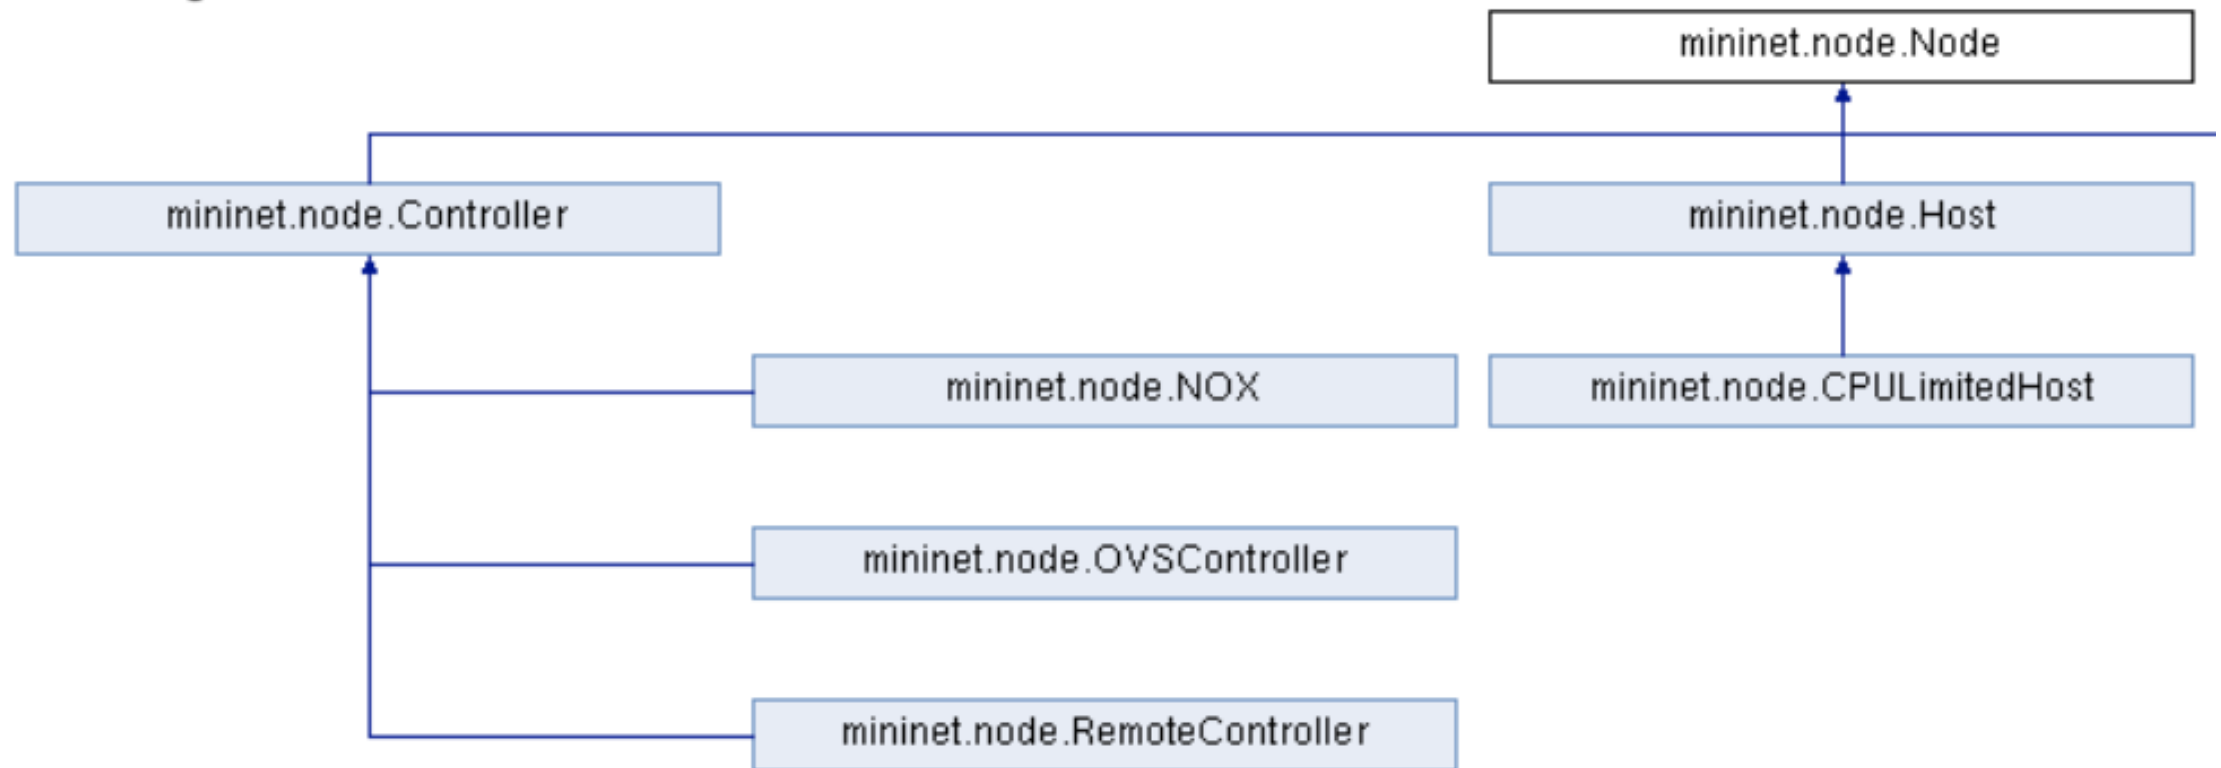
```
c0.start()
s1.start( [ c0 ] )
print h1.cmd( 'ping -c1', h2.IP() )
s1.stop()
c0.stop()
```

Node class and subclasses (1/2)



- ▶ Node generic class
- ▶ 3 subclasses: **Controller**, **Host**, **Switch**

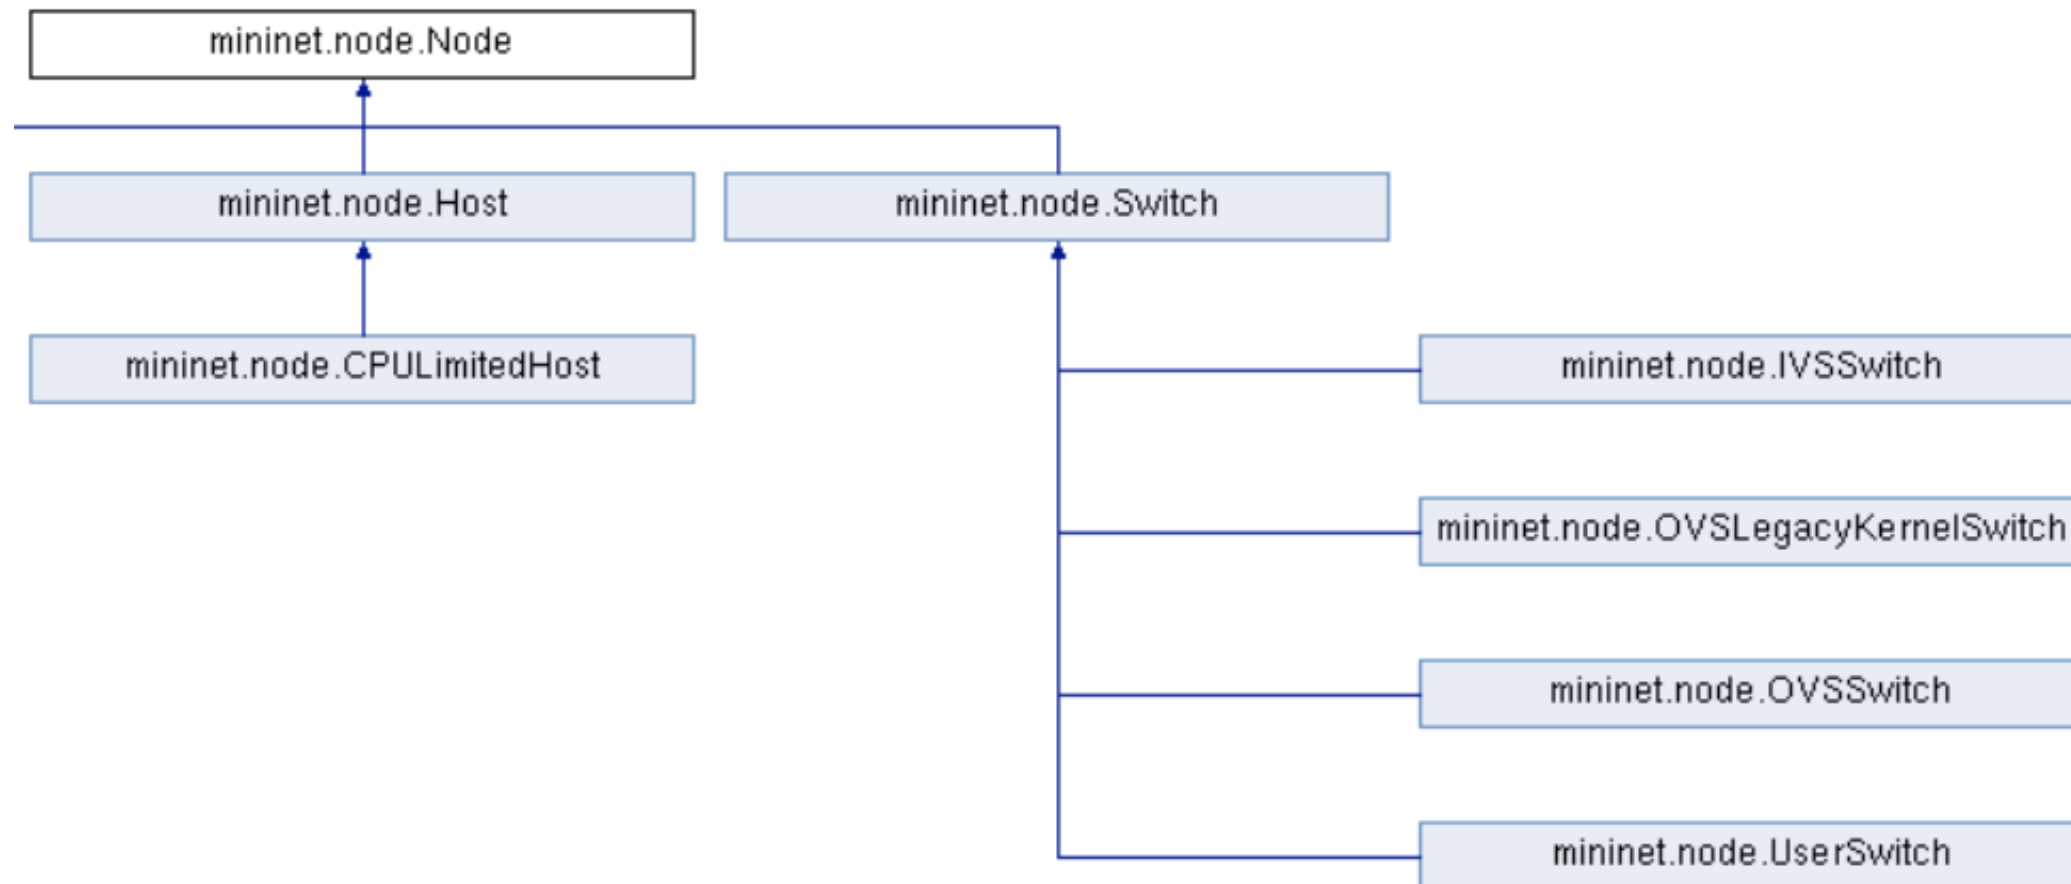
Inheritance diagram for mininet.node.Node:



Node class and subclasses (2/2)



- ▶ Node generic class
- ▶ 3 subclasses: Controller, Host, Switch



► mininet.net.Mininet

► Network emulation with hosts spawned in network namespaces

Class	Method	Description
Mininet	addHost	Add a host to network
	addSwitch	Add a switch to network
	addLink	Link two nodes into together
	addController	Add a controller to network
	getNodeByName	Return node(s) with given name(s)
	start	Start controller and switches
	stop	Stop the controller, switches and hosts
	ping	Ping between all specified hosts and return all data

```
net = Mininet()  
h1 = net.addHost( 'h1' )  
h2 = net.addHost( 'h2' )  
s1 = net.addSwitch( 's1' )  
c0 = net.addController( 'c0' )  
net.addLink( h1, s1 )  
net.addLink( h2, s1 )
```

```
net.start()  
print h1.cmd( 'ping -c1', h2.IP() )  
CLI( net )  
net.stop()
```

► mininet.topo.Topo

Class	Method	Description
Topo	Methods similar to net	E.g., addHost, addSwitch, addLink,
	addNode	Add node to graph
	addPort	Generate port mapping for new edge
	switches	Return all switches
	Hosts/nodes/switches/links	Return all hosts/nodes/switches/links
	isSwitch	Return true if node is a switch, return false otherwise

```
class SingleSwitchTopo( Topo ):
```

```
    "Single Switch Topology"
```

```
    def build( self, count=1):
```

```
        hosts = [ self.addHost( 'h%d' % i )
```

```
            for i in range( 1, count + 1 ) ]
```

```
        s1 = self.addSwitch( 's1' )
```

```
        for h in hosts:
```

```
            self.addLink( h, s1 )
```

```
net = Mininet( topo=SingleSwitchTopo(3) )
```

```
net.start()
```

```
CLI( net )
```

```
net.stop()
```

First Mininet script (1/2)



```
#!/usr/bin/python
from mininet.topo import Topo
from mininet.net import Mininet
from mininet.util import dumpNodeConnections
from mininet.log import setLogLevel

class SingleSwitchTopo(Topo):
    "Single switch connected to n hosts."
    def build(self, n=2):
        switch = self.addSwitch('s1')
        # Python's range(N) generates 0..N-1
        for h in range(n):
            host = self.addHost('h%s' % (h+1))
            self.addLink(host, switch)

#...
```

Custom topology
class

First Mininet script (2/2)



```
# ...  
def simpleTest():  
    "Create and test a simple network"  
    topo = SingleSwitchTopo(n=4)  
    net = Mininet(topo)  
    net.start()  
    print("Dumping host connections")  
    dumpNodeConnections(net.hosts)  
    print("Testing network connectivity")  
    net.pingAll()  
    net.stop()  
  
if __name__ == '__main__':  
    # Tell mininet to print useful information  
    setLogLevel('info')  
    simpleTest()
```

Custom test function

Script startup

Second Mininet script (1/2)



```
#!/usr/bin/python
from mininet.topo import Topo
from mininet.net import Mininet
from mininet.node import CPULimitedHost
from mininet.link import TCLink
from mininet.util import dumpNodeConnections
from mininet.log import setLogLevel

class SingleSwitchTopo(Topo):
    "Single switch connected to n hosts."
    def build(self, n=2):
        switch = self.addSwitch('s1')
        for h in range(n):
            # Each host gets 50%/n of system CPU
            host = self.addHost('h%s' % (h+1), cpu=.5/n)
            # 10 Mbps, 5ms delay, 2% loss, 1000 packet queue
            self.addLink(host, switch, bw=10, delay='5ms',
                          loss=2, max_queue_size=1000,
                          use_htb=True)
#...
```

Custom topology class

Second Mininet script (2/2)



```
# ...  
def perfTest():  
    "Create network and run a simple performance test"  
    topo = SingleSwitchTopo(n=4)  
    net = Mininet(topo)  
    net.start()  
    print("Dumping host connections")  
    dumpNodeConnections(net.hosts)  
    print("Testing network connectivity")  
    net.pingAll()  
    # print("Testing bandwidth between h1 and h4")  
    # h1, h4 = net.get('h1', 'h4')  
    # net.iperf( (h1, h4) )  
    net.stop()  
  
if __name__ == '__main__':  
    setLogLevel('info')  
    perfTest()
```

Custom test function

Script startup

Third Mininet script (1/3)



```
#!/usr/bin/python
from mininet.topo import Topo
from mininet.net import Mininet
from mininet.node import CPULimitedHost
from mininet.link import TCLink
from mininet.util import dumpNodeConnections
from mininet.log import setLogLevel

class LinearTopo(Topo):
    "Linear topology of n switches, with one host per switch."
    def build(self, n=2):
        lastSwitch = None
        for i in range(n):
            # Each host gets 50%/n of system CPU
            host = self.addHost('h%s' % (i+1), cpu=.5/n)
            switch = self.addSwitch('s%s' % (i+1))
            # 10 Mbps, 5ms delay, 2% loss, 1000 packet queue
            self.addLink(host, switch, bw=10, delay='5ms',
                          loss=2, max_queue_size=1000, use_htb=True)
        #...
```

Custom topology
class

Third Mininet script (2/3)



```
#...
if lastSwitch:
    self.addLink(switch, lastSwitch, bw=10, delay='5ms',
                 loss=2, max_queue_size=1000, use_htb=True)
    lastSwitch = switch

def perfTest():
    "Create network and run a simple performance test"
    topo = SingleSwitchTopo(n=4)
    net = Mininet(topo)
    net.start()
    print("Dumping host connections")
    dumpNodeConnections(net.hosts)
    print("Testing network connectivity")
    net.pingAll()
    # print("Testing bandwidth between h1 and h4")
    # h1, h4 = net.get('h1', 'h4')
    # net.iperf( (h1, h4) )
    net.stop()

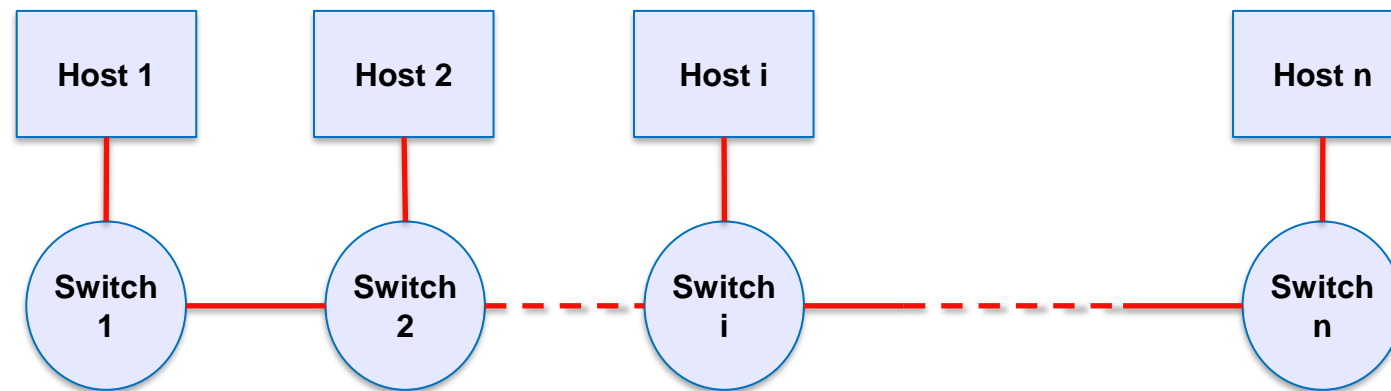
#...
```

Custom test function

```
#...  
if __name__ == '__main__':  
    setLogLevel('info')  
    perfTest()
```

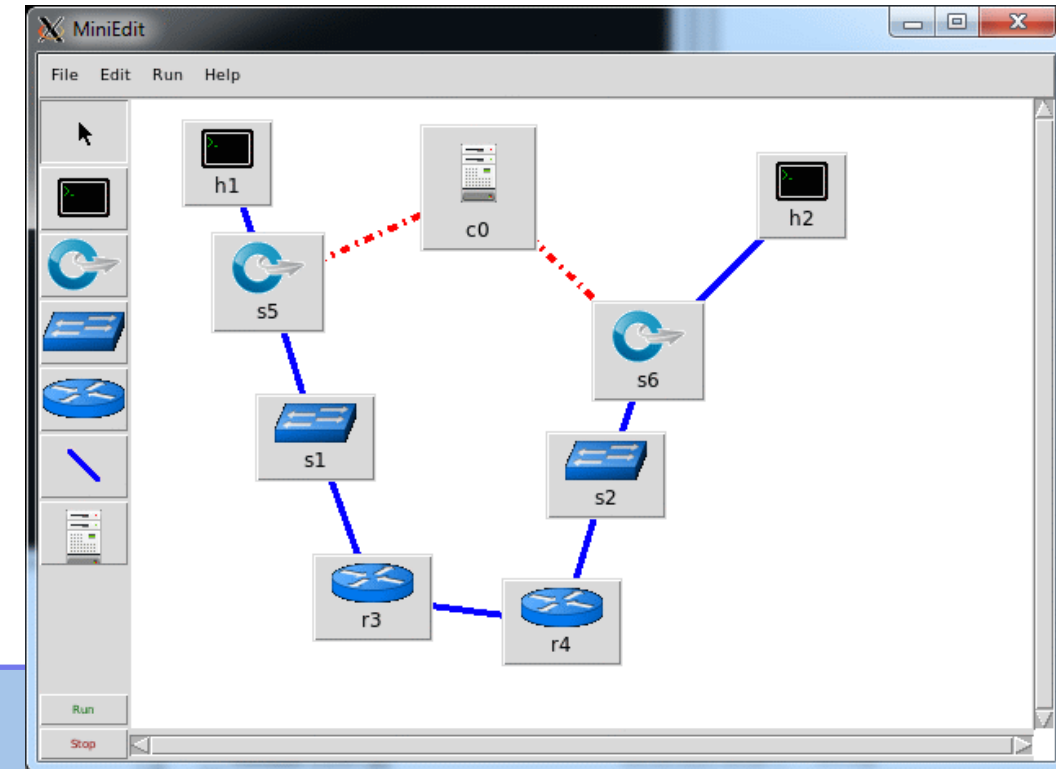
Script startup

Linear topology



► MiniEdit

- A GUI application which eases the Mininet topology generation
- Either save the topology or export as a Mininet python script



```
$ sudo apt install python-tk
$ sudo apt-get install openvswitch-testcontroller
$ sudo ln /usr/bin/ovs-testcontroller /usr/bin/controller
$ sudo python2 ~/mininet/examples/miniedit.py

# create the topology

# export the topology as a .py file

$ sudo python3 my_topology.py
```

Remote controller with Ryu



- ▶ Ryu is a component-based software defined networking framework.
- ▶ Ryu provides software components with well defined API that make it easy for developers to create new network management and control applications.
- ▶ We use Ryu for running an example of remote controller.

```
$ cd /home/mininet/.local/lib/python3.10/site-packages/ryu/app
$ nano wsgi.py
# Comment the AlreadyHandledResponse class in wsgi.py
$ ryu-manager simple_switch_13.py
```

```
$ sudo mn --controller remote
```

```
$ sudo wireshark
```

Controller logic: Simple Switch 13



```
from ryu.base import app_manager
from ryu.controller import ofp_event
from ryu.controller.handler import CONFIG_DISPATCHER, MAIN_DISPATCHER
from ryu.controller.handler import set_ev_cls
from ryu.ofproto import ofproto_v1_3
from ryu.lib.packet import packet
from ryu.lib.packet import ethernet
from ryu.lib.packet import ether_types
```

```
class SimpleSwitch13(app_manager.RyuApp):
```

```
    OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]
```

```
    def __init__(self, *args, **kwargs):
```

```
        super(SimpleSwitch13, self).__init__(*args, **kwargs)
```

```
        self.mac_to_port = {}
```

```
    @set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
```

```
    def switch_features_handler(self, ev):
```

```
    def add_flow(self, datapath, priority, match, actions, buffer_id=None):
```

```
    @set_ev_cls(ofp_event.EventOFPPacketIn, MAIN_DISPATCHER)
```

```
    def _packet_in_handler(self, ev):
```

Controller logic: Simple Switch 13



```
def switch_features_handler(self, ev):
    datapath = ev.msg.datapath
    ofproto = datapath.ofproto
    parser = datapath.ofproto_parser
    # install table-miss flow entry
    #
    # We specify NO BUFFER to max_len of the output action due to
    # OVS bug. At this moment, if we specify a lesser number, e.g.,
    # 128, OVS will send Packet-In with invalid buffer_id and
    # truncated packet data. In that case, we cannot output packets
    # correctly. The bug has been fixed in OVS v2.1.0.
    match = parser.OFPMatch()
    actions = [parser.OFPActionOutput(ofproto.OFPP_CONTROLLER, ofproto.OFPCML_NO_BUFFER)]
    self.add_flow(datapath, 0, match, actions)

def add_flow(self, datapath, priority, match, actions, buffer_id=None):
    ofproto = datapath.ofproto
    parser = datapath.ofproto_parser
    inst = [parser.OFPInstructionActions(ofproto.OFPIT_APPLY_ACTIONS, actions)]
    if buffer_id:
        mod = parser.OFPFlowMod(datapath=datapath, buffer_id=buffer_id,
                                priority=priority, match=match,
                                instructions=inst)
    else:
        mod = parser.OFPFlowMod(datapath=datapath, priority=priority,
                                match=match, instructions=inst)
    datapath.send_msg(mod)
```

Controller logic: Simple Switch 13



```
@set_ev_cls(ofp_event.EventOFPPacketIn, MAIN_DISPATCHER)
def _packet_in_handler(self, ev):
    # If you hit this you might want to increase
    # the "miss_send_length" of your switch
    if ev.msg.msg_len < ev.msg.total_len:
        self.logger.debug("packet truncated: only %s of %s bytes",
                           ev.msg.msg_len, ev.msg.total_len)

    msg = ev.msg
    datapath = msg.datapath
    ofproto = datapath.ofproto
    parser = datapath.ofproto_parser
    in_port = msg.match['in_port']
    pkt = packet.Packet(msg.data)
    eth = pkt.get_protocols(ethernet.ethernet)[0]
    if eth.ethertype == ether_types.ETH_TYPE_LLDP:
        # ignore lldp packet
        return
    dst = eth.dst
    src = eth.src
    dpid = datapath.id
    self.mac_to_port.setdefault(dpid, {})
    self.logger.info("packet in %s %s %s %s", dpid, src, dst, in_port)

    # learn a mac address to avoid FLOOD next time.
    self.mac_to_port[dpid][src] = in_port

    if dst in self.mac_to_port[dpid]:
        out_port = self.mac_to_port[dpid][dst]
    else:
        out_port = ofproto.OFPP_FLOOD

    actions = [parser.OFPACTIONOutput(out_port)]

    # install a flow to avoid packet_in next time
    if out_port != ofproto.OFPP_FLOOD:
        match = parser.OFPMATCH(in_port=in_port, eth_dst=dst, eth_src=src)
        # verify if we have a valid buffer_id, if yes avoid to send both
        # flow_mod & packet_out
        if msg.buffer_id != ofproto.OFP_NO_BUFFER:
            self.add_flow(datapath, 1, match, actions, msg.buffer_id)
            return
        else:
            self.add_flow(datapath, 1, match, actions)
    data = None

    if msg.buffer_id == ofproto.OFP_NO_BUFFER:
        data = msg.data
    out = parser.OFPPACKETOut(datapath=datapath, buffer_id=msg.buffer_id,
                              in_port=in_port, actions=actions, data=data)
    datapath.send_msg(out)
```

Example: topology defined in Python script and remote controller



```
#!/usr/bin/python
import threading
import random
import time
from mininet.log import setLogLevel, info
from mininet.topo import Topo
from mininet.net import Mininet, CLI
from mininet.node import OVSKernelSwitch, Host
from mininet.link import TCLink, Link
from mininet.node import RemoteController #Controller

...
```


Example: topology defined in Python script and remote controller



```
class Environment(object):
    def __init__(self):
        "Create a network."
        self.net = Mininet(controller=RemoteController, link=TCLink)
        info("*** Starting controller\n")
        c1 = self.net.addController( 'c1', controller=RemoteController) #Controller
        c1.start()
        info("*** Adding hosts and switches\n")
        self.h1 = self.net.addHost('h1', mac='00:00:00:00:00:01', ip='10.0.0.1')
        self.h2 = self.net.addHost('h2', mac='00:00:00:00:00:02', ip='10.0.0.2')
        self.cpe1 = self.net.addSwitch('s1', cls=OVSKernelSwitch)
        self.cpe2 = self.net.addSwitch('s2', cls=OVSKernelSwitch)
        self.core1 = self.net.addSwitch('s3', cls=OVSKernelSwitch)
        info("*** Adding links\n")
        self.net.addLink(self.h1, self.cpe1, bw=6, delay='0.0025ms')
        self.path1 = self.net.addLink(self.cpe1, self.core1, bw=3, delay='25ms')
        self.net.addLink(self.core1, self.cpe2, bw=3, delay='25ms')
        self.net.addLink(self.cpe2, self.h2, bw=6, delay='0.0025ms')
        info("*** Starting network\n")
        self.net.build()
        self.net.start()
    ...
```



```
...  
if __name__ == '__main__':  
  
    setLogLevel('info')  
    info('starting the environment\n')  
    env = Environment()  
  
    info("*** Running CLI\n")  
    CLI(env.net)
```

Running the example



```
$ ryu-manager simple_switch_13.py
```

```
$ sudo python3 my_topo.py
```

```
$ sudo wireshark
```

Capture OpenFlow packets with wireshark



- ▶ If wireshark is not able to decode OF packets, reinstall a newer version

```
sudo apt-get remove wireshark
sudo apt-get -y install libgtk-3-dev libqt4-dev flex bison
wget https://www.wireshark.org/download/src/all-versions/wireshark-1.12.3.tar.bz2
tar xvfj wireshark-1.12.3.tar.bz2
cd wireshark-1.12.3
./configure
make -j4
sudo make install
sudo echo "/usr/local/lib" >> /etc/ld.so.conf
sudo ldconfig
```

- ▶ If the controller is running locally, capture packets on **lo** interface (*loopback*) on port TCP/6653 (filter = tcp port 6653)